

Assignment Project Exam Help

Databases: Introduction

<https://eduassistpro.github.io>

Imperial College London

Add WeChat edu_assist_pr

Databases are Computer Stores of Data!

Tiny Bank Ltd Customer: McBrien, P.
Strand Branch Current Acc: 10000100
Sortcode: 55-66-67

Trans	Amount	Date
1000	2300.00	5/1/1999
1002	1232.41	3/1/1999
1006	10.23	15/1/1999

Tiny Bank Ltd Customer: Poulouvassilis, A.
Wimbledon Branch Current Acc: 10000107
Sortcode: 55-66-56

Trans	Amount	Date
1004	-116.00	11/1/1999
1007	315.56	15/1/1999

Tiny Bank Ltd Customer: [redacted]
Strand Branch Current Acc: [redacted]
Sortcode: 5

Trans	Amount	Date
1001	4000.00	5/1/1999
1008	1230.00	15/1/1999

1009	5600.00	18/1/1999
------	---------	-----------

Tiny Bank Ltd Customer: Boyd, M.
Goodge St Branch Current Acc: 10000103
Sortcode: 55-66-34

Trans	Amount	Date
1005	145.50	12/1/1999

Ti
Wi
So

Trans	Amount	Date
-------	--------	------

No transactions this month

Deposit Rates

Account	Rate
101	5.25
119	5.50

Relational Data Model

Relational Data Model

Roughly: storing data in tables

no	sortc	branch	tdate
100		00.00	1999-01-05
101		00.00	1999-01-05
100	67 Strand	34005.00 current	McBrien, P. 1002 -223.45 1999-01-08
107	56 Wimbledon	84340.45 current	Poulov 1999-01-11
103	34 Goode St	6900.67 current	Boyd, M 1999-01-12
100	67 Strand	34005.00 current	McBrien 1999-01-15
107	56 Wimbledon	84340.45 current	Poulov 1999-01-15
101	67 Strand	34005.00 deposit	McBrien 1999-01-15
119	56 Wimbledon	84340.45 deposit	Poulovassilis, A. 5.50 1009 5600.00 1999-01-18


<https://eduassistpro.github.io/>

Add WeChat [edu_assist_pro](https://eduassistpro.github.io/)

Database Design: ER Modelling

Assignment Project Exam Help

<https://eduassistpro.github.io>



Structured Data: Relational Model

sortcode	branch	cash
56	Wimbledon	94340.45
34	'Goodge St'	8900.67
67		

no	type	name	rate?	sortcode
100	'current'	'McBrien, P.'	NULL	67
101	'deposit'	'McBrien, P.'	5.25	67
			LL	34
			LL	56
			50	56
			-	50

mid	no		
1000	100		
1001	101	4000.00	5/1/1999
1002	100	-223.45	8/1/1999
1004	107	-100.00	11/1/1999
1005	103	145.50	12/1/1999
1006	100	11.23	15/1/1999
1007	107	345.56	15/1/1999
1008	101	1230.00	15/1/1999
1009	119	5600.00	18/1/1999

key branch(so

key branch(b

key movement

key account(n

movement(n \Rightarrow account(sortcode) \xRightarrow{fk} branch(sortcode)

Data Model: CSV

Assignment Project Exam Help

branch.csv

sort code, bname, cash

56,"Wimbledon",94340.45

34,

67,

movement.csv

mid, no, amount, tdate

1000,100,2300.00,5/1/1999

999

999

1999

999

99

<https://eduassistpro.github.io>

no, type, c

100,"current", "McBrien, P.", ,67

101,"deposit", "McBrien, P.", 5.25,67

103,"current", "Bord, M.", ,35

107,"current", "Poulovassilis, A.", ,56

119,"deposit", "Poulovassilis, A.", 5.50,56

125,"current", "Bailey, J.", ,56

1007,107,345.56,15/1/1999

Add WeChat edu_assist_pro

Semistructured Data: XML

Assignment Project Exam Help

```
<bank>
  <branch sortcode="67" bname="Strand" cash="34005.00" >
    <account
      <move
        <move
          </account>
        <account
          <move
            <movement mid="1008" amount="1230.00" tdate="15/1/1999" />
          </account>
        </branch>
      </bank>
```

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

SQL DDL: Implementation of the Relational Model

```
CREATE TABLE branch
( sortcode INTEGER NOT NULL,
  bname VARCHAR(20) NOT NULL,
  cash DECIMAL(10,2) NOT NULL,
  CONSTRAINT branch_pk PRIMARY KEY (sortcode)
)
```

```
CREATE UNI
ON branch(b
```

```
CREATE TABLE account
( no INTEGER NOT NULL,
  type CHAR(8) NOT NULL,
  cname VARCHAR(20) NOT NULL,
  rate DECIMAL(4,2) NOT NULL,
  sortcode INTEGER NOT NULL,
  CONSTRAINT account_pk
    PRIMARY KEY (no),
```

ENCES branch

N account(type)

<https://eduassistpro.github.io>

Add WeChat [edu_assist_pro](https://eduassistpro)

```
CREA
(mid IN
no IN
amo
tdate DATETIME NOT NULL,
CONSTRAINT movement_pk
    PRIMARY KEY (mid),
CONSTRAINT movement_fk
    FOREIGN KEY (no) REFERENCES account
)
```


SQL DML: Implementation of the Relational Algebra

Basic SQL SELECT statements

```
SELECT no, cname, rate
FROM account
WHERE type='deposit'
```

SQL Joins

```
SELECT b
FROM b
WHERE type='deposit'
```

Same as

```
SELECT bname, no, rate
FROM account JOIN branch ON branch.so
WHERE type='deposit'
```

Same as

```
SELECT bname, no, rate
FROM account, branch
WHERE branch.sortcode=account.sortcode
AND type='deposit'
```

RDBMS Products

Assignment Project Exam Help

Product	SQL Language	Company
---------	--------------	---------

<https://eduassistpro.github.io>

PostgreSQL	PL/pgSQL	Open Source
MySQL	MySQL	

Add WeChat edu_assist_pro

All partially implement ANSI SQL

Transactions

```
BEGIN TRANSACTION
```

```
UPDATE branch
SET cash = cash - 100.00
WHERE sortcode=56
```

database management systems

```
UP
```

```
SE
```

```
W
```

```
COMMIT TRANSACTION
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

The ACID Properties

- **Atomicity** all or nothing
- **Consistency** consistent before → consistent after
- **Isolation** independent of any other transaction
- **Durability** completed transaction are durable

Add WeChat edu_assist_pro

Transaction Properties: Atomicity

```
BEGIN TRANSACTION
```

```
UPDATE branch
```

```
SET cash=cash-1000.00
```

```
WHERE sortcode=56
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

Failure to m

Suppose that the system crashes half way through processing the first part of the transfer has been written to disc

- The database on disc is left in an inconsistent state: the sum £137,246.12 but only £127,246.12 recorded
- A DBMS implementing **Atomicity** of transactions would on restart undo the change to branch 56

Transaction Properties: Consistency

```
BEGIN TRANSACTION
DELETE FROM branch
WHERE sortcode = 56
INSERT INTO account
VALUES
END TRAN
```

<https://eduassistpro.github.io>

Failure to maintain consistency

Suppose that a user deletes branch with sortcode 56, and inserts a new account number 100 for John Smith at branch sortcode 34.

- The database is left in an inconsistent state for two reasons:
 - it has three accounts recorded for a branch that appears not to exist, and
 - it has two records for account number 100, with different details for the account
- A DBMS implementing **Consistency** of transactions would forbid both of these changes to the database

Transaction Properties: Isolation

BEGIN TRANSACTION

UPDATE branch

SET cash=cash-10000.00

WHERE sortcode=56

BEGIN TRANSACTION

SELECT SUM(cash) AS net_cash

UPDA

SET

WHERE sortcode=34

END TRANSACTION

END

Failure to maintain Isolation

Suppose that the system sums the cash in the bank in one transaction, half way through processing a cash transfer in another transaction

- The result of the summation of cash in the bank erroneously reports £127,246.12, whereas the movement of cash always leaves a total of £137,246.12
- A DBMS implementing **Isolation** of transactions ensures that transactions always report results based on the values of committed transactions

Transaction Properties: Durability

```
BEGIN TRANSACTION
```

```
UPDATE branch
```

```
SET cash=cash -10000.00
```

```
WHERE brn_code=56
```

```
UPDATE branch
```

```
SET
```

```
WHE
```

```
END TRAN
```

```
CRASH
```

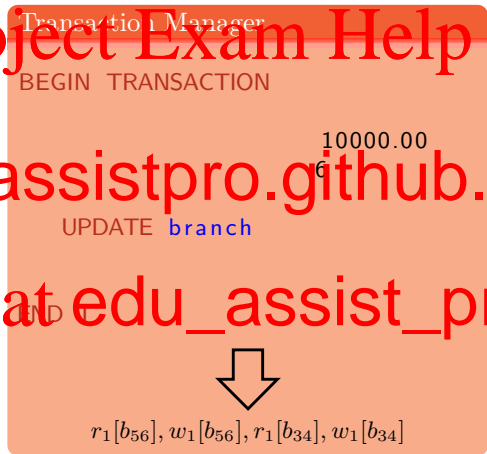
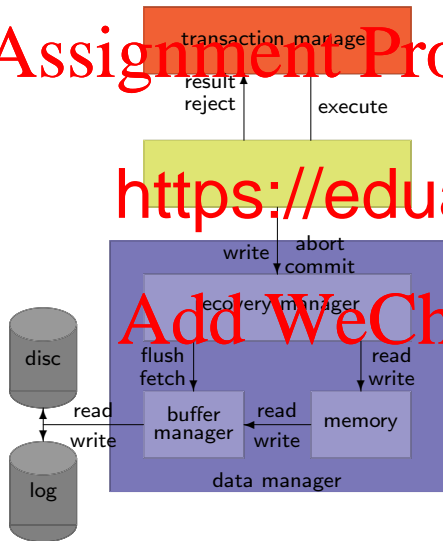
<https://eduassistpro.github.io>

Failure to maintain Durability

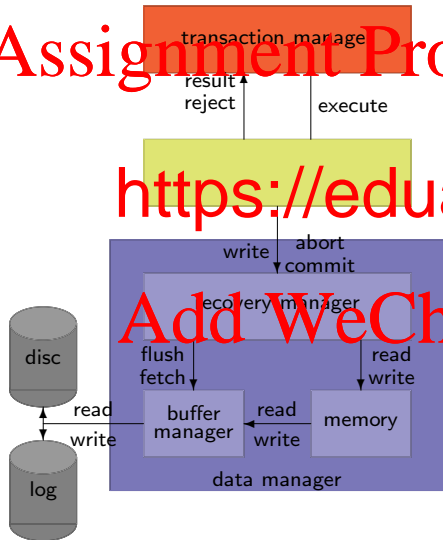
Suppose that the system crashes after informing the user the transfer of cash, but has not yet written to disc the update to

- The database on disc is left in an inconsistent state, with £10,000 'missing'
- A DBMS implementing **Durability** of transactions would on restart complete the change to branch 34 (or alternatively never inform a user of commitment with writing the results to disc).

DBMS Architecture



DBMS Architecture



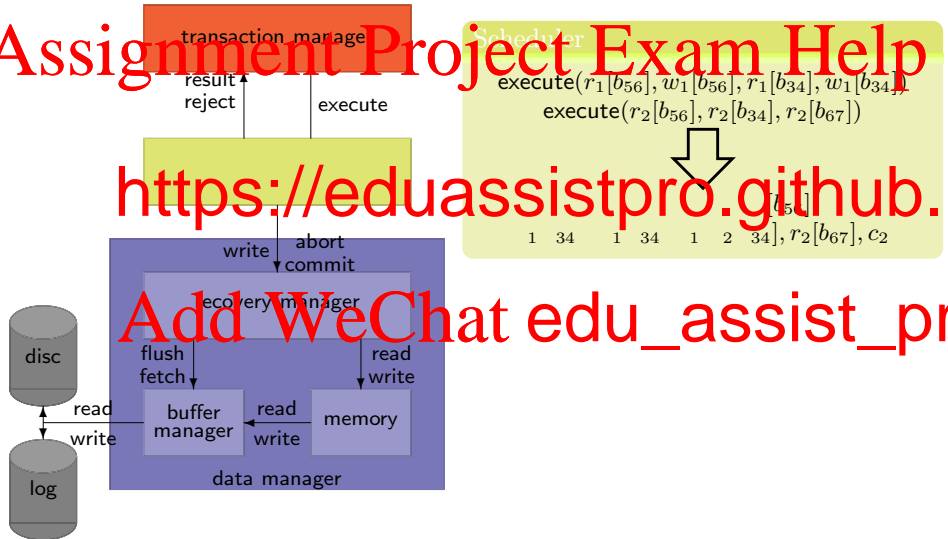
Transaction Manager

```
BEGIN TRANSACTION
AS net_cash
```

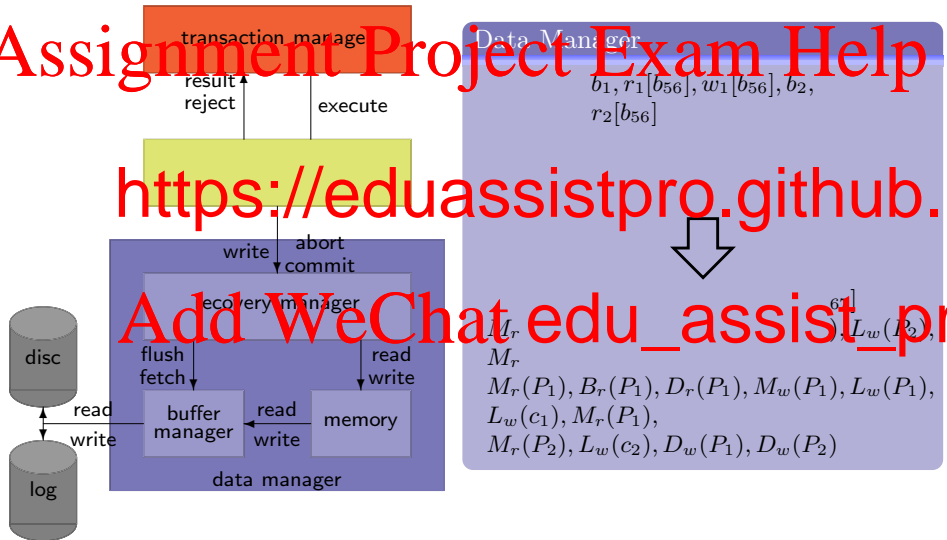
<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

DBMS Architecture




DBMS Architecture



ANSI/SPARC Model

Assignment Project Exam Help



<https://eduassistpro.github.io>

Add WeChat [edu_assist_pro](https://eduassistpro.github.io)


- ANSI/SPARC model views three levels of abstractions
- **schema** means *structure of the database*

ANSI/SPARC Model (Internal Schema)

Assignment Project Exam Help

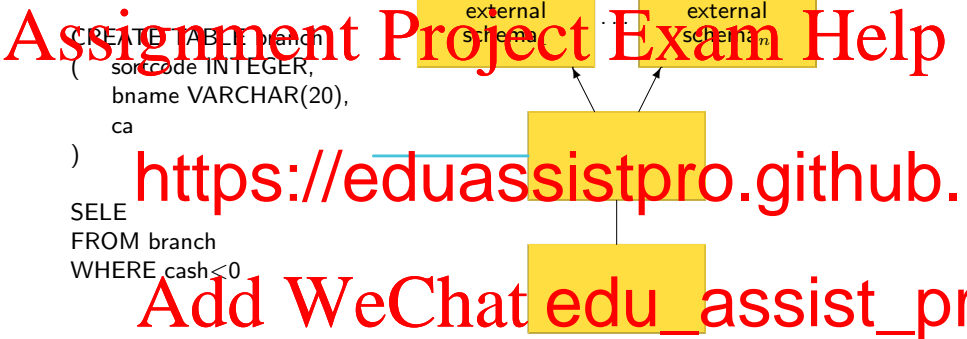
<https://eduassistpro.github.io>

- 2k page size
- B-tree index
- Strings end with char 0



- Describes the physical layout of data

ANSI/SPARC Model (Conceptual Schema)



- defined in **data definition language (DDL)**
- queried using **data manipulation language (DML)**
- controlled by **database administrator (DBA)**

ANSI/SPARC Model (External Schema)

```
CREATE VIEW bust
SELECT src code
FROM branch
WHERE cash < 0
```

external
schema₁

...

external
schema_n

<https://eduassistpro.github.io>

Add WeChat [edu_assist_pr](#)

- Define a schema for a particular user/application

Course Format

Schedule

- Three hours combined lectures/tutorials per week, running into week 10
- Coursework that helps you prepare for the exam
- May Exam

Books

Several good books, some of which are available in more advanced courses are:

- *Fundamentals of Database Systems*, 6th Ed, Elmasri and Navathe, Addison-Wesley
- *Database Systems: The Complete Book*, 2nd Ed, Garcia-Molina, Ullman and Widom, Pearson
- *Database Systems*, 5th Ed, Connolly and Begg, Addison Wesley

Course Resources

Course Web Site

<https://www.doc.ic.ac.uk/~pjm/hb/>

- Lecture slides
- Exam

■ <https://eduassistpro.github.io>

Resources

- **CATe** course work handout and submission
- **Piazza** discussion forum
- **email** course email list

If you are not on Level 2 on CATe, nothing works!

Course Content

Conceptual Layer: Relational Algebra

- SQL
- Datalog

Conceptu

- Prope
- Database design using ER models

Physical Layer: Transaction Processing

- Serialisability
- Recovery and Checkpointing